

Billy Herzberg

william.herzberg@marquette.edu

Building Convolutional Neural Networks with TensorFlow 2.x



MARQUETTE
UNIVERSITY

BE THE
DIFFERENCE.



About Me


- BS Mechanical Engineering 2018 – Marquette University
- MS Applied Statistics 2020 – Marquette University
- PhD Computational Sciences 20?? – Marquette University

- First started with deep learning in Summer 2019
- Convolutional Neural Networks for Electrical Impedance Tomography (EIT) image reconstruction
 - Post processing images with a U-Net CNN
 - Model-based, learned, iterative reconstruction

What Makes a Neural Network?

- Data Set
 - Inputs X
 - Truths Y
 - Split into training and validation
 - $X = [X_{tr} \quad X_{va}]^T$
- Model
 - Has layers
 - Has parameters θ
 - Makes predictions
 - $\hat{Y}^{(k)} = Net_{\theta}(X^{(k)})$

- Loss Function
 - Evaluates model
 - $loss = f(Y, \hat{Y})$
- Optimizer
 - Adjusts θ during training
 - Ex. gradient descent
- Training the model
 - Split training data into batches
 - Update θ after each batch



Outline


- 1D Convolution
 - Stride, Padding
- 2D Convolution
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channel output
- Choosing Weights
 - Finding changes, edges, peaks, features, etc.
- Pooling Layers
- Transposed Convolutions
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- U-Net Architecture
 - Post-processing images

Outline


- **1D Convolution**
 - Stride, Padding
- **2D Convolution**
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channel output
- **Choosing Weights**
 - Finding changes, edges, peaks, features, etc.
- **Pooling Layers**
- **Transposed Convolutions**
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- **U-Net Architecture**
 - Post-processing images

1D Convolutions


Input →



Kernel →




Output →




1D Convolutions

Input →




Kernel →




$$(5 * 3) + (4 * 1) + (3 * 2) = 25$$

Output →




1D Convolutions

Input →




Kernel →




$$(4 * 3) + (3 * 1) + (8 * 2) = 31$$

Output →




1D Convolutions

Input →




Kernel →



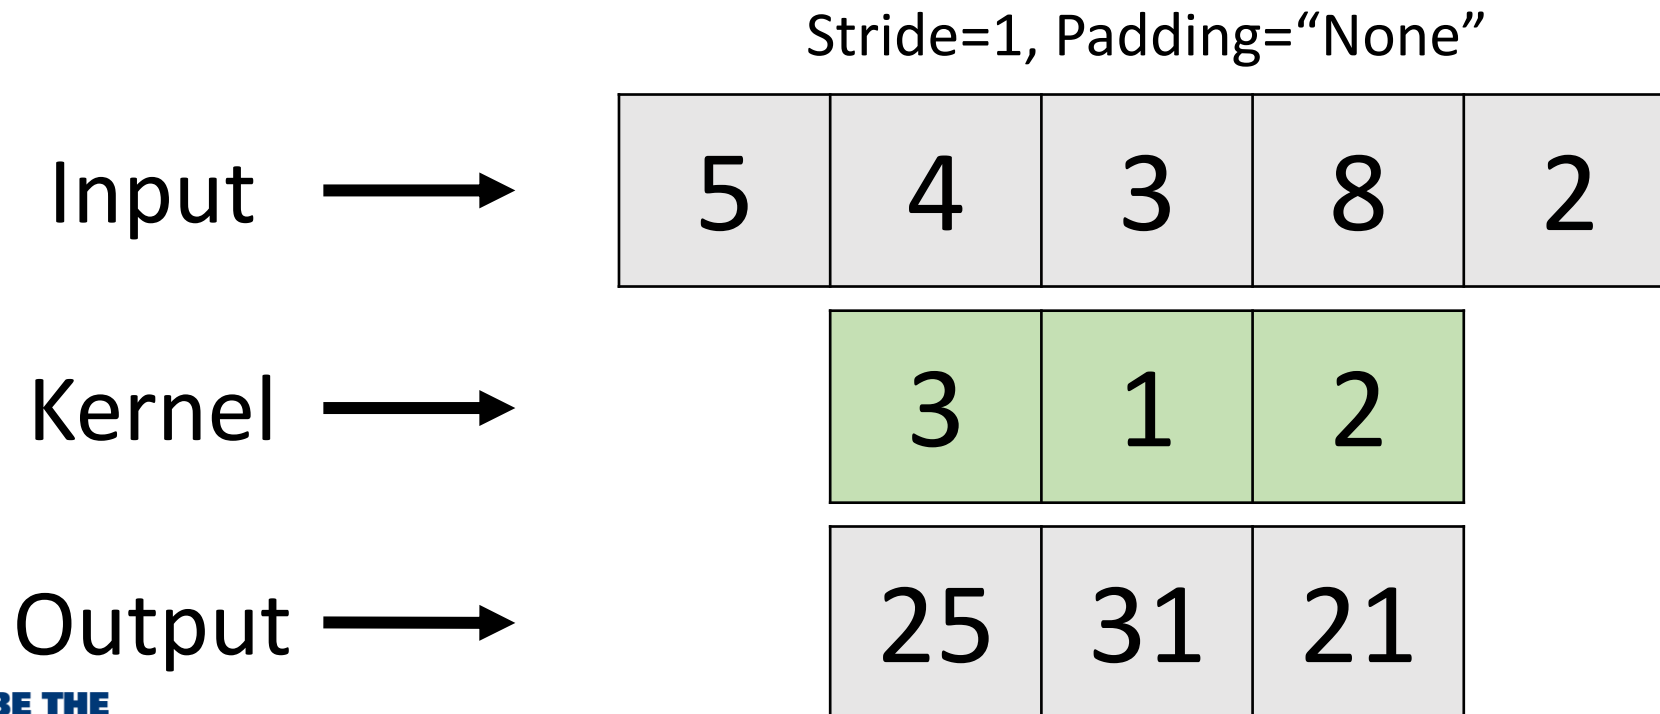
$$(3 * 3) + (8 * 1) + (2 * 2) = 21$$

Output →



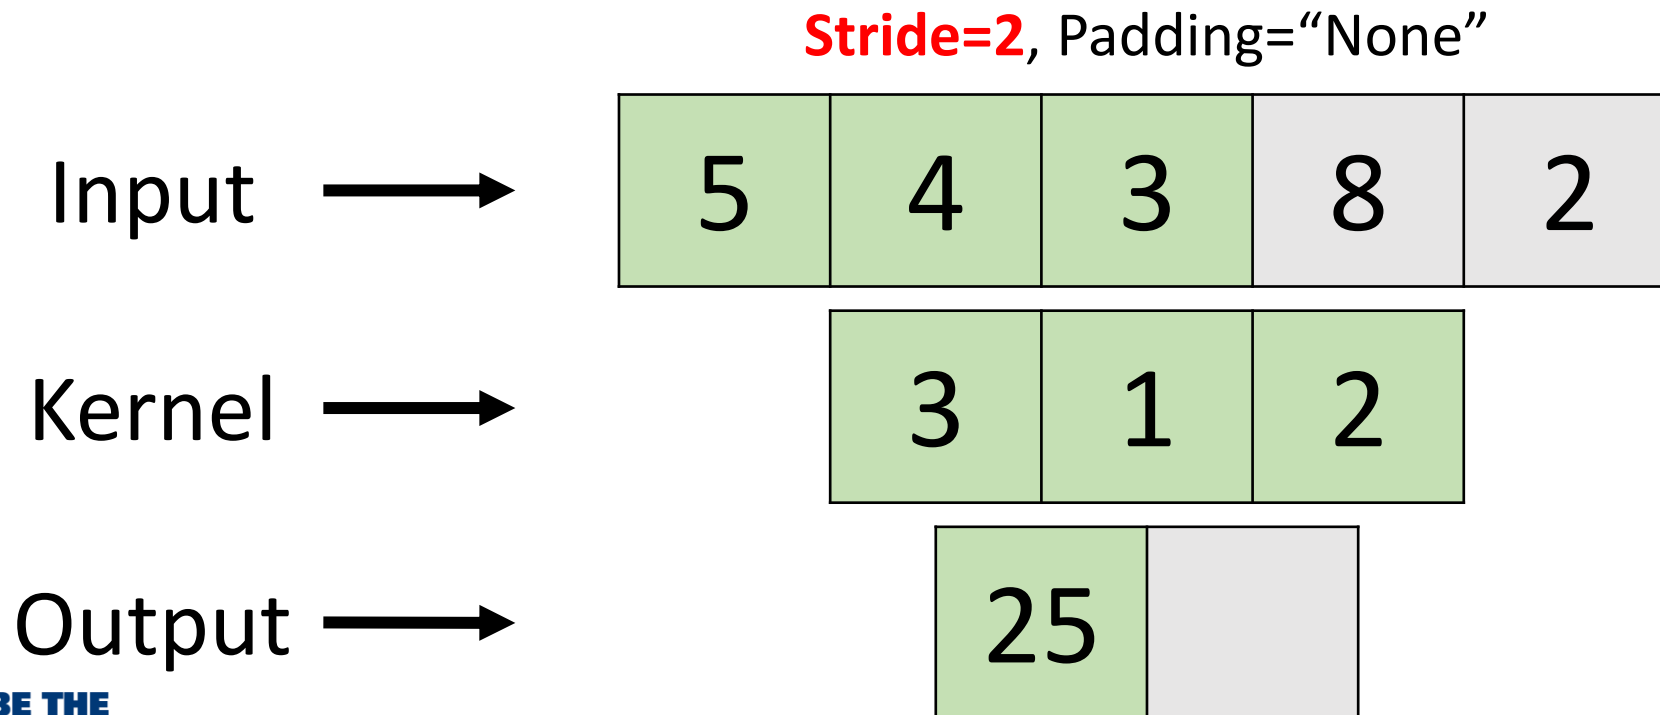
1D Convolutions

- **Stride** is the distance the kernel moves for each output calculation
- **Padding** can be used to get the desired size for the output




1D Convolutions

- **Stride** is the distance the kernel moves for each output calculation
- Padding can be used to get the desired size for the output



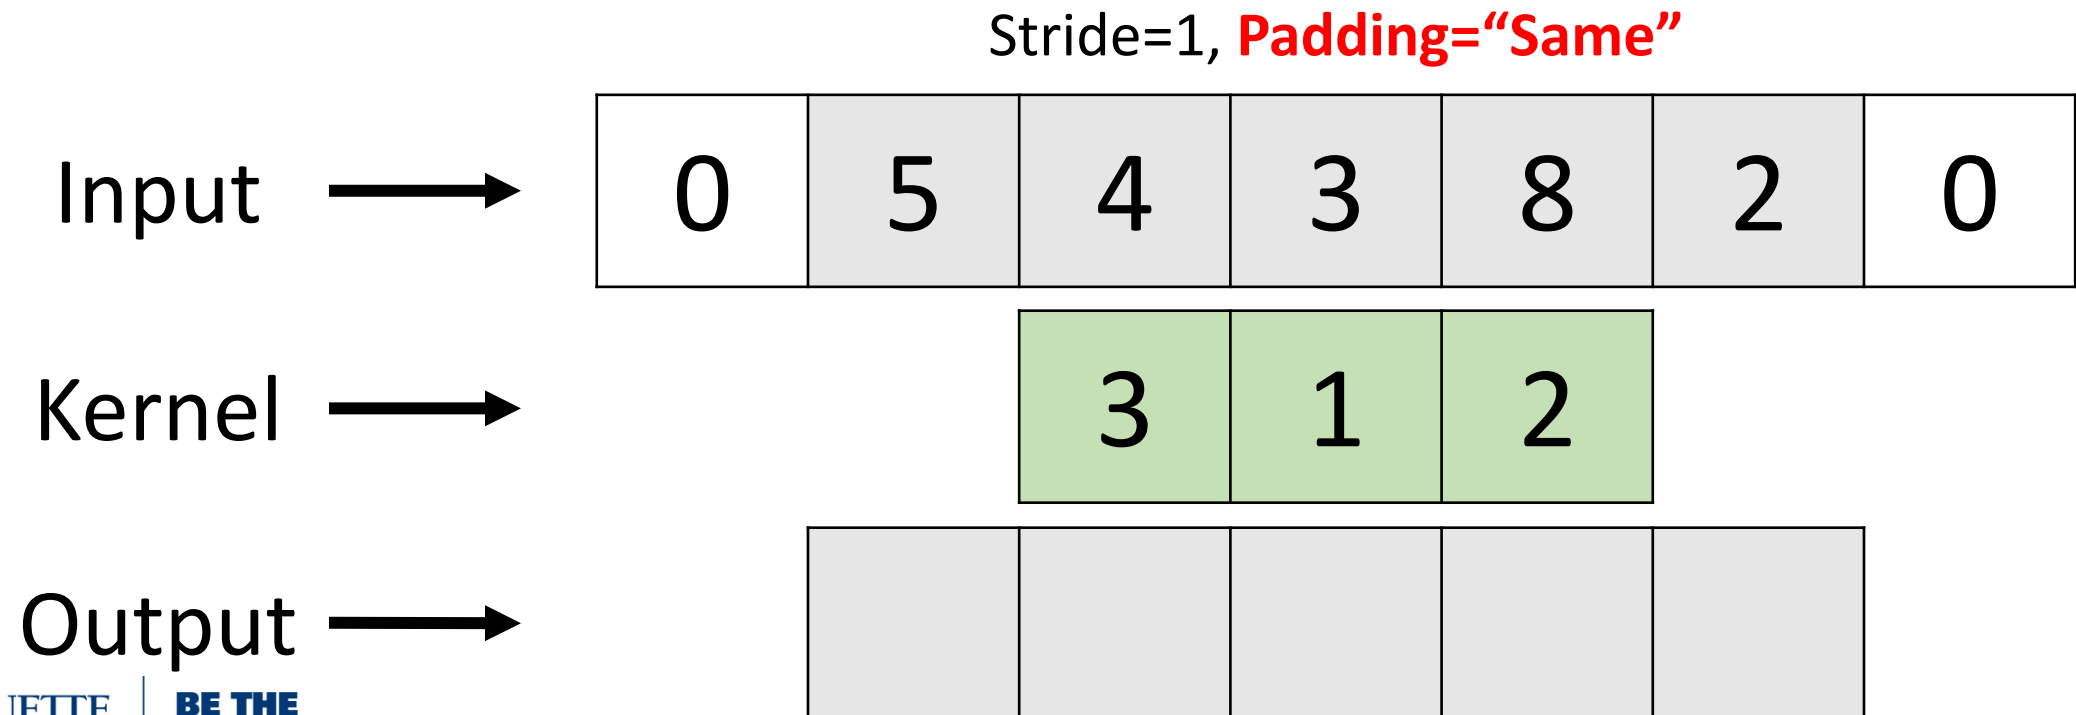
1D Convolutions

- **Stride** is the distance the kernel moves for each output calculation
- Padding can be used to get the desired size for the output



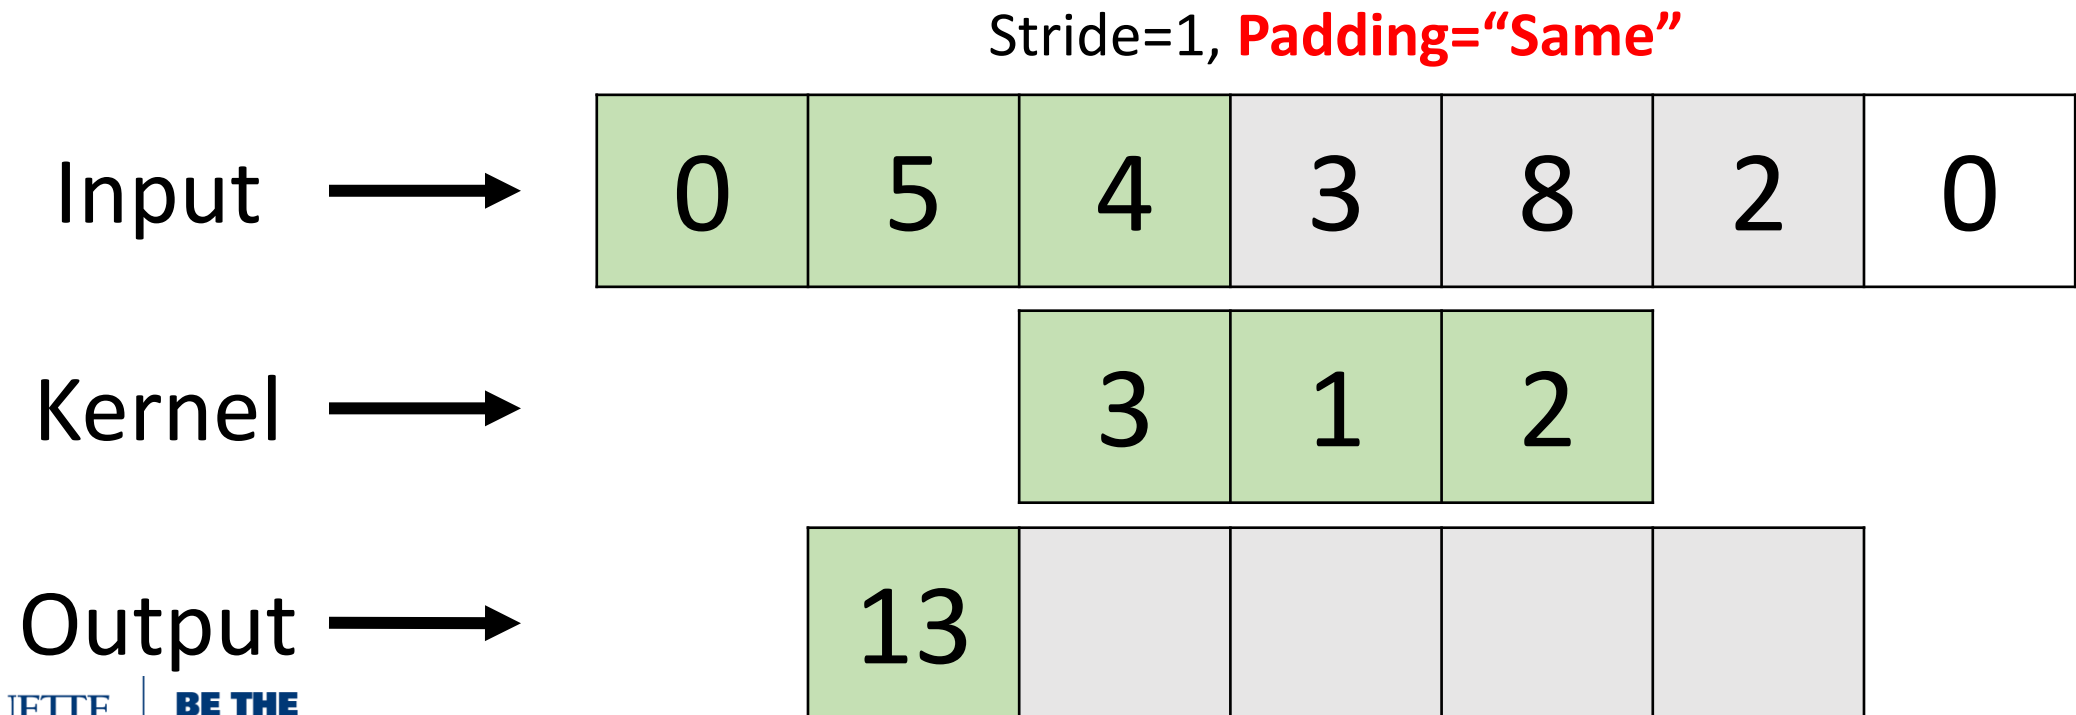
1D Convolutions

- Stride is the distance the kernel moves for each output calculation
- **Padding** can be used to get the desired size for the output



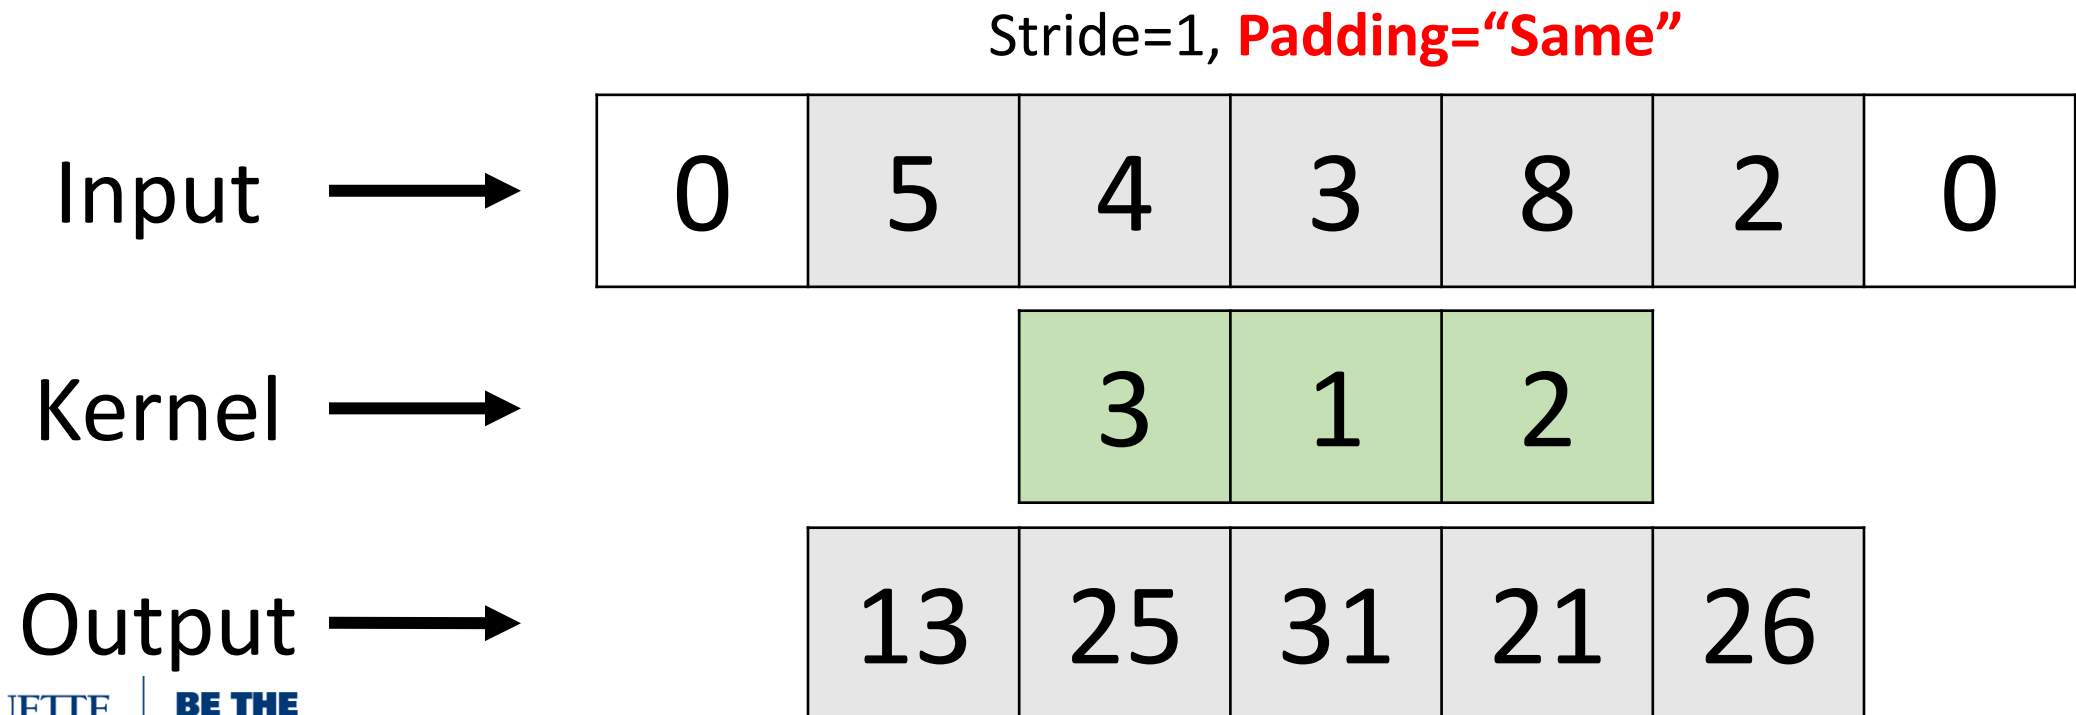
1D Convolutions

- Stride is the distance the kernel moves for each output calculation
- **Padding** can be used to get the desired size for the output



1D Convolutions

- Stride is the distance the kernel moves for each output calculation
- **Padding** can be used to get the desired size for the output



Questions?

- [illegible]

Outline

- 1D Convolution
 - Stride, Padding
- **2D Convolution**
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channel output
- Choosing Weights
 - Finding changes, edges, peaks, features, etc.
- Pooling Layers
- Transposed Convolutions
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- U-Net Architecture
 - Post-processing images

2D Convolutions

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

2D Convolutions

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

18		

$$\begin{aligned} &(5 * 1) + (4 * 0) + (3 * 1) + \dots \\ &(8 * 0) + (1 * 2) + (2 * 1) + \dots \\ &(1 * 0) + (7 * 0) + (3 * 2) = 18 \end{aligned}$$

2D Convolutions

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

18	24	

$$\begin{aligned} &(4 * 1) + (3 * 0) + (8 * 1) + \dots \\ &(1 * 0) + (2 * 2) + (4 * 1) + \dots \\ &(7 * 0) + (3 * 0) + (2 * 2) = 24 \end{aligned}$$

2D Convolutions

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

18	24	26
27	14	21
26	14	19

$$\begin{aligned} &(3 * 1) + (2 * 0) + (6 * 1) + \dots \\ &(0 * 0) + (1 * 2) + (4 * 1) + \dots \\ &(7 * 0) + (3 * 0) + (2 * 2) = 19 \end{aligned}$$

2D Convolutions

- Again, stride and padding can be used to manipulate output size

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

Stride=(1,1)
Padding="None"

18	24	26
27	14	21
26	14	19

2D Convolutions

- Again, stride and padding can be used to manipulate output size

Input

5	4	3	8	2
8	1	2	4	1
1	7	3	2	6
2	4	0	1	4
3	1	7	3	2

Kernel

1	0	1
0	2	1
0	0	2

Output

Stride=(2,2)
Padding="None"

18	26
26	19

2D Convolutions

- Again, stride and padding can be used to manipulate output size

Input

0	0	0	0	0	0	0
0	5	4	3	8	2	0
0	8	1	2	4	1	0
0	1	7	3	2	6	0
0	2	4	0	1	4	0
0	3	1	7	3	2	0
0	0	0	0	0	0	0

Kernel

1	0	1
0	2	1
0	0	2


Output

Stride=(1,1)
Padding="Same"

16	15	22	20	4
35	18	24	26	10
18	27	14	21	16
17	26	14	19	10
11	11	22	12	5


2D Convolutions

- The dimension of the convolution is determined by the stride dimension, NOT the input or output dimension
- RGB images have three channels (one each for Red, Green, and Blue) and are therefore 3D arrays. Still, 2D convolutions are used




2D Convolutions

Input




Kernel




Output

Stride=(1,1)
Padding="None"




2D Convolutions

Input




Kernel



Output


Stride=(1,1)
Padding="None"




Kernel must have same number of channels as input

2D Convolutions


Input



Kernel



Output



Stride=(1,1)
Padding="None"

Kernel must have same number of channels as input

Number of kernels determines number of output channels

2D Convolutions

2	3	1
1	1	2
3	3	2

1	0	1
0	2	1
0	0	2

2	1	3
1	2	1
1	0	0

Input

2	3	1		
1	1	2		
3	3	2		

2D Convolutions

2	3	1
1	1	2
3	3	2

1	1
0	2

1	0	1
0	2	1
0	0	2

1	0
1	1

2	1	3
1	2	1
1	0	0

1	1
2	2

Input

2	3	1
1	1	2
3	3	2

Kernel

1	1
0	2

2D Convolutions

2	3	1
1	1	2
3	3	2

1	1
0	2

7	

1	0	1
0	2	1
0	0	2

1	0
1	1

2	1	3
1	2	1
1	0	0

1	1
2	2

Input

2	3	1
1	1	2
3	3	2

Kernel

1	1
0	2

2D Convolutions

2	3	1
1	1	2
3	3	2

1	1
0	2

7	8
8	7

1	0	1
0	2	1
0	0	2

1	0
1	1

2	1	3
1	2	1
1	0	0

1	1
2	2

Input

2	3	1
1	1	2
3	3	2

Kernel

1	1
0	2

2D Convolutions

2	3	1
1	1	2
3	3	2

1	1
0	2

7	8
8	7

1	0	1
0	2	1
0	0	2

1	0
1	1

3	3
0	4

2	1	3
1	2	1
1	0	0

1	1
2	2

9	10
5	3

Input

2	3	1
1	1	2
3	3	2

Kernel

1	1
0	2

2D Convolutions

2	3	1
1	1	2
3	3	2

1	0	1
0	2	1
0	0	2

2	1	3
1	2	1
1	0	0

1	1
0	2

1	0
1	1

1	1
2	2

7	8
8	7

+

3	3
0	4

+

9	10
5	3

=

19	21
13	14

Output

Input


2	3	1
1	1	2
3	3	2

Kernel


1	1
0	2

2D Convolutions

Input




Kernel




Output

Stride=(1,1)
Padding="None"




2D Convolutions

Input




Kernel



Output

Stride=(1,1)
Padding="None"



Number of kernels
determines number
of output channels

Questions?


- 1D Convolutions
 - 2D Convolutions
 - Multiple input channels
 - Multiple output channels
-
- Up next: Choosing weights

Outline

- 1D Convolution
 - Stride, Padding
- 2D Convolution
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channels on output
- **Choosing Weights**
 - Finding changes, edges, peaks, features, etc.
- Pooling Layers
- Transposed Convolutions
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- U-Net Architecture
 - Post-processing images

Choosing Weights

- Can we choose a kernel so that the output is the change of the input?



Choosing Weights

- We can! Selecting $(w_1, w_2) = (-1, 1)$ will output the change

Input

0	0	1	0	2	1	2	0
---	---	---	---	---	---	---	---

Kernel

-1	1
----	---

Output

0	1	-1	2	-1	1	-2
---	---	----	---	----	---	----

Choosing Weights

- Other kernels can be used to detect **peaks**, valleys, features, etc.
- Same idea works with multidimensional kernels
- Note that padding and stride=1 were used here

Input

0	0	1	0	2	1	2	0	0
---	---	---	---	---	---	---	---	---

Kernel

-1	2	-1
----	---	----

Output

-1	2	-3	3	-2	3	-2
----	---	----	---	----	---	----

Choosing Weights

- So what kernels do we pick for our CNN?
- Luckily, we don't need to pick, the kernels are learned during training!
- For each convolutional layer, how many parameters need to be learned by the algorithm?

$$(num\ elements\ in\ kernel) * (num\ kernels) + (num\ kernels)$$

weights

biases

Questions?

- 1D Convolutions
 - 2D Convolutions
 - Multiple Channels
 - Choosing weights to highlight features
-
- Up next: Pooling layers and a classification example

Outline

- 1D Convolution
 - Stride, Padding
- 2D Convolution
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channels on output
- Choosing Weights
 - Finding changes, edges, peaks, features, etc.
- **Pooling Layers**
- Transposed Convolutions
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- U-Net Architecture
 - Post-processing images


Pooling Layers

- Pooling layers reduce the resolution (height and/or width)
- Extract dominant features that are positionally invariant
- Reduce noise
- Slide a window over the input and take the max (or mean, etc.) of the input values as the output
- Max pooling is typically used in CNN's because it disregards noise better than other options

Max Pooling Layers

Input

1	2	4	1
7	3	2	6
4	0	1	4
1	7	3	2




Output

7		

Max Pooling Layers

Input

1	2	4	1
7	3	2	6
4	0	1	4
1	7	3	2




Output

7	4	

Max Pooling Layers

Input

1	2	4	1
7	3	2	6
4	0	1	4
1	7	3	2




Output

7	4	6

Max Pooling Layers

Input

1	2	4	1
7	3	2	6
4	0	1	4
1	7	3	2



Output


7	4	6
7	3	6
7	7	4

Max Pooling Layers

- Pooling layers are defined by kernel size and stride
- There are no learned parameters in a pooling layer

Input

1	2	4	1
7	3	2	6
4	0	1	4
1	7	3	2



Output

7	4	6
7	3	6
7	7	4



Kernel Size=(2,2)
Stride=(1,1)

Questions?

- 1D Convolutions
 - 2D Convolutions
 - Multiple Channels
 - Choosing weights to highlight features
 - Max Pooling
-
- Up next: Classification Example in Google Colab

Classification Example

- Network Architecture
 - Input Layer
 - Convolutional Layer
 - 1 kernel
 - Pooling Layer
 - Output Layer
- Sigmoid Activation
- Binary Cross Entropy



Outline


- 1D Convolution
 - Stride, Padding
- 2D Convolution
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channels on output
- Choosing Weights
 - Finding changes, edges, peaks, features, etc.
- Pooling Layers
- **Transposed Convolutions**
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- U-Net Architecture
 - Post-processing images

Transposed Convolutions


- Also called backwards convolution or deconvolution among other things (transposed makes the most sense in my opinion)
- A form of up sampling
- Increases the resolution (height and width) of the input
 - Kernel size, stride, and padding help manage output sizes
- Usually used while decreasing the number of channels

Transposed Convolutions (1D)

Input →




Kernel →




$$(2 * 3) = 6 \quad (2 * 2) = 4$$

Output →




Transposed Convolutions (1D)

Input →




Kernel →




$$(1 * 3) = 3 \quad (1 * 2) = 2$$

Output →




Transposed Convolutions (1D)

Input →




Kernel →



$$(4 * 3) = 12 \quad (4 * 2) = 8$$

Output →



Transposed Convolutions (1D)

- This example used a stride of 1 and no padding

Input →

2	1	4
---	---	---

Kernel →

3	2
---	---


Output →

6	7	14	8
---	---	----	---


Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries

Input →



Kernel →




Output →




Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries

Input →



Kernel →




Output →




Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries


Input →



Kernel →




Output →




Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries


Input →



Kernel →




Output →




Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries


Input →



Kernel →




Output →




Transposed Convolutions (1D)

- We can pad with zeros NOT on the ends, but between entries


Input →



Kernel →



Output →



Transposed Convolutions (1D)

- Padding the input with 0's has the same effect as adding a stride to the output

Input →

2	0	1	0	4
---	---	---	---	---

Kernel →

3	2
---	---


Output →

6	4	3	2	12	8
---	---	---	---	----	---

Transposed Convolutions (1D)

- Padding the input with 0's has the same effect as adding a stride to the output (stride=2)


Input →



Kernel →



Output →



Questions?

- 1D Convolutions
 - 2D Convolutions
 - Multiple Channels
 - Choosing weights to highlight features
 - 1D Transposed Convolutions
-
- Up next: 2D Transposed Convolutions

Transposed Convolutions (2D)

- Again, 2D refers to the dimension of the stride parameter
- Padding between input entries with zeros or adjusting the stride of the output work with 2D transposed convolutions too

Transposed Convolutions (2D)

- Example with no padding (output stride = (1,1))

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	1	
1	0	

Transposed Convolutions (2D)

- Example with no padding (output stride = (1,1))

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	1+2	2
1	0+2	0

Transposed Convolutions (2D)

- Example with no padding (output stride = (1,1))

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	1+2	2
1+6	0+2+3	0
3	0	

Transposed Convolutions (2D)

- Example with no padding (output stride = (1,1))

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	1+2	2
1+6	0+2+ 3+8	0+4
3	0+4	0

Transposed Convolutions (2D)

- Example with no padding (output stride = (1,1))

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	3	2
7	13	4
3	4	0

Transposed Convolutions (2D)

- The stride of (1,1) resulted in overlap in the orange output entries

Input

1	2
3	4

Kernel

2	1
1	0

Output

2	3	2
7	13	4
3	4	0

Transposed Convolutions (2D)

These mean the same thing.
It's not two separate things
being done; it's two ways of
thinking of the same thing

- Example with padding (output stride = (2,2))

Input

1	0	2
0	0	0
3	0	4

Kernel

2	1
1	0

Output

2	1	4	2
1	0	2	0
6	3	8	4
3	0	4	0

Transposed Convolutions (2D)

Input

3	1
-1	2

Kernel

1	-1
2	1

Output

3	-3		
6	3		

1	1
2	1

1	0
-1	-1

Input

3	1	
-1	2	

Kernel

1	-1	
2	1	

Transposed Convolutions (2D)

Input

3	1
-1	2

Kernel

1	-1
2	1

Output

3	-3	1	-1
6	3	2	1
-1	1	2	-2
-2	-1	4	2

+

1	1
2	1

1	0
-1	-1

1	0	1	0
-1	-1	-1	-1
2	0	1	0
-2	-2	-1	-1

=

Input

3	1	
-1	2	

Kernel


1	-1	
2	1	

Output



4	-3	2	-1
5	2	1	0
1	1	3	-2
-4	-3	3	1

Transposed Convolutions (2D)

Input




Kernel



Kernels must have same
number of channels as input

Output



Number of kernels
determines number
of output channels

Transposed Convolution Layer

- For each transposed convolutional layer, how many parameters need to be learned by the algorithm?

$$\underbrace{(\textit{num elements in kernel}) * (\textit{num kernels})}_{\text{weights}} + \underbrace{(\textit{num kernels})}_{\text{biases}}$$

Questions?


- 1D Convolutions
 - 2D Convolutions
 - Multiple Channels
 - Choosing weights to highlight features
 - 1D Transposed Convolutions
 - 2D Transposed Convolutions
-
- Up next: U-Net Architecture for post processing images

Outline

- 1D Convolution
 - Stride, Padding
- 2D Convolution
 - Single channel input
 - Multiple channel input (i.e. RGB)
 - Multiple channels on output
- Choosing Weights
 - Finding changes, edges, peaks, features, etc.
- Pooling Layers
- Transposed Convolutions
 - 1D
 - 2D
 - Stride, padding
 - Multiple channel input
 - Multiple channel output
- **U-Net Architecture**
 - Post-processing images



U-Net Architecture

- Presented by Ronneberger et al. in 2015 for image segmentation
- Shown to sharpen EIT D-bar reconstructions by Hamilton and Hauptmann in 2018





U-Net Architecture

- Contracting path (left) identifies features
- Expanding path (right) offers localization



U-Net Architecture

- Contracting path (left) identifies features
- Expanding path (right) offers localization




Questions?

- 1D Convolutions
- 2D Convolutions
- Multiple Channels
- Choosing weights to highlight features
- 1D Transposed Convolutions
- 2D Transposed Convolutions
- U-Net Architecture
- Up next: U-Net Example in Google Colab

U-Net Example

- Input is a blurry image with (or without) a smiley face
- Truth is an image with a ring (or without) and no noise

- Network
 - U-Net
 - 1 Max Pool Layer
 - MSE loss function
 - ReLU activation



Why use CNNs? (as opposed to dense networks)

- Number of weights is not related to the image sizes
 - Could have used 500x500 smiley face images
- Feature learning
 - By the end of training, the classifier knew what features to look for
- Spatially invariant
 - Can detect a smiley face in any location

Thank You!

Questions?